

MethylScore Manual

version 0.1.13+

Computomics GmbH
Eisenbahnstr. 1 · 72072 Tübingen · Germany

+4970715683995
info@computomics.com
<https://computomics.com>

Management

Dr. Sebastian J. Schultheiss

Scientific Advisory Board

Prof. Dr. Detlef Weigel
Prof. Dr. Daniel Huson
Prof. Dr. Gunnar Rätsch
Prof. Dr. Karsten Borgwardt

Strategic Advisory Board

Dr. Ralf Allrutz
Reinhard Hamburger

Legal

Computomics GmbH, Tübingen, Germany
Management Dr. Sebastian J. Schultheiss
Commercial Register AG Stuttgart HRB742810
EU VAT ID DE286347612

Contents

MethylScore strategy	4
#1.1 MR calling strategy	4
#1.2 DMR calling strategy	6
 Running MethylScore	 8
#2.1 Prerequisites	8
#2.2 Folder structure	9
#2.3 Input	9
#2.4 Output file formats	10
#2.5 Performance recommendations	13
#2.6 Command line parameters	13
Quick Usage	13
#2.7 License	17

This document describes the strategy and usage of MethylScore, a pipeline to call differential methylation between samples using next-generation sequencing.

MethylScore strategy

The pipeline consists of several modules that can be run as a whole or one by one. The overall workflow and the used external tools with their versions are outlined in Figure 1.

In brief, the pipeline starts from mapping files of each sample or of each technical replicate. Mapping files are expected to be in bam format, produced by the bisulfite read mapping tool `bwa-meth`¹ [Pedersen et al. (2014)].

First, the mapping files of technical replicates are merged, each sample's mappings are de-duplicated using `picardtools`, then split by chromosome, and for each sample and chromosome, the numbers of methylated/unmethylated reads per position (pileup information) are retrieved using the tool `MethylExtract`. The pileup information from all analyzed samples is summarized in a so-called genome matrix that is generated per sample and chromosome in parallel.

The global genome matrix serves as input for the detection of methylated regions per sample and the conversion to an `igv` file that can be used to visualize methylation information in the Integrative Genomics Viewer (IGV). Methylated regions (MRs) are determined by a two-state Hidden Markov Model (HMM)-based method that learns different methylation level distributions for an unmethylated and a methylated state from whole genome data (described in section #1.1).

Finally, to obtain significant differences in methylation on a regional scale between different samples, MethylScore, in short, clusters samples by methylation levels and statistically tests the group's methylation distributions for significant differences (described in section #1.2).

#1.1 MR calling strategy

MR calling is performed by applying a Hidden Markov Model that fits context-specific beta binomial distributions from whole genome data for an unmethylated and a methylated state. Given the trained model, sites are probabilistically classified into the most likely state via Posterior Decoding (default) or Viterbi's algorithm. The basic implementation was adapted from Molaro et al. (2011).

Desert size parameter

The advantage of using an HMM is that it takes the directly preceding state into account. Chromosomes are 'natural' split points where the HMM decoding starts over again without previous state information. The HMM implementation described here has additional split points defined by a maximum genomic distance between adjacent covered cytosines. Regions of absent read coverage or cytosines are referred to as 'deserts' in the original implementation. In practice, it means that an MR cannot span a region without methylation information of the size of a desert (MethylScore parameter `DESERT_SIZE`, default: 100 bp).

¹ <https://github.com/brentp/bwa-meth/archive/master.zip> (last accessed June 2017)

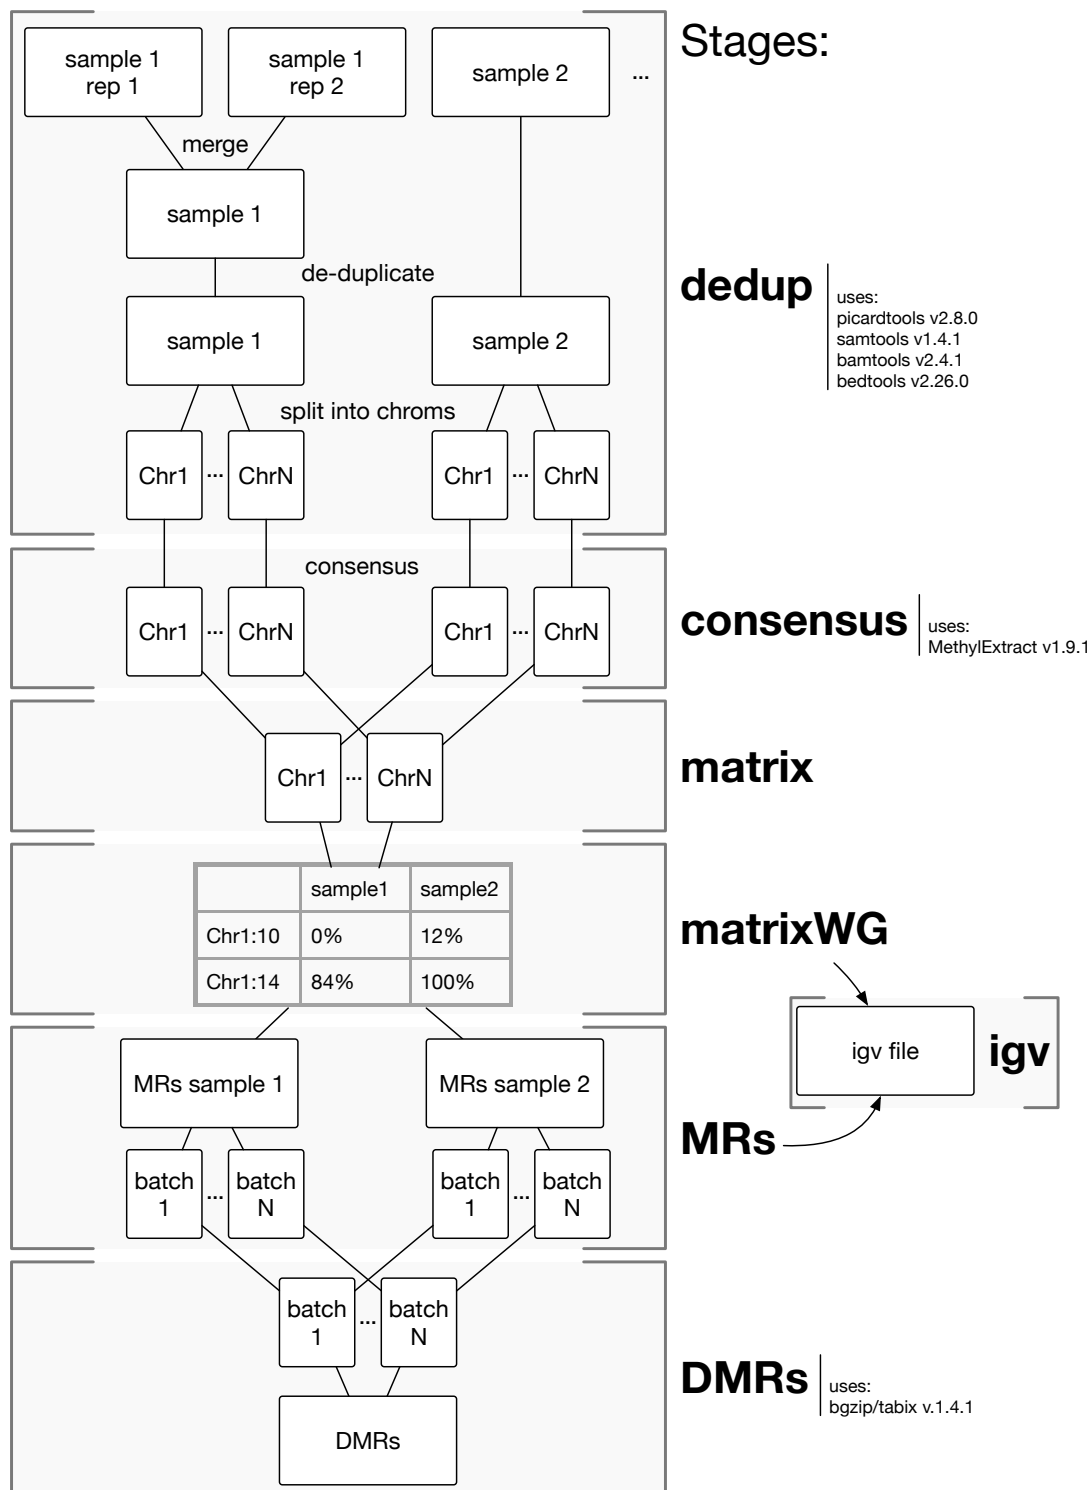


Figure 1: MethylScore analysis pipeline. An example with two samples is shown, where sample 1 has two technical replicates (rep1, rep2). Horizontally arranged boxes represent tasks that can be run in parallel.

Post-processing parameters

Since consecutive sites in methylated state ("high methylation regions") can be very short, there are two alternative ways to finally define MRs:

1. Permutation test: all covered cytosines are randomly shuffled and 'random' high methylation regions are called. High methylation regions are scored by the sum of the contained methylation rates. The scores of the 'real' high methylation regions are tested against the 'random' score distribution to calculate p-values. After FDR calculation, high methylation regions with an FDR of 0.01 (default) are defined as MRs.
This strategy proved reasonable on whole-genome libraries and for species where most genome space is unmethylated (e.g. plants from the *Arabidopsis* genus and rice). This strategy is followed if the MethylScore parameter `MR_MIN_C` is set to 0.
2. Minimum number of cytosines: high methylation regions containing a minimum number of covered cytosines are defined as MRs.
This simple length filter proves useful for methylation-enriched libraries or for species with a high fraction of methylated sequence, when a random shuffling of the genome already produces regions of moderate to high methylation density by chance. Then, moderately dense regions will be discarded by the permutation test. This second strategy remedies the described disadvantage of the permutation test and is chosen if the MethylScore parameter `MR_MIN_C` is set to a positive value.

Besides an MR length filter as just described above, there are two additional post-processing filters to facilitate a biologically meaningful interpretation. Nearby MRs within a certain base pair distance can be merged into one region (MethylScore parameter `MERGE_DIST`, default: 30 bp, turned off: 0 bp), and lowly methylated positions can be trimmed off the ends of MRs (MethylScore parameter `TRIM_METHRATE`, default: sites below 10 % methylation rate are trimmed off, turned off: 0 %).

#1.2 DMR calling strategy

For calling DMRs, MethylScore focuses on the unified methylated genome space determined by the sample-specific MRs. To reduce complexity and ease the computation of a large number of samples, MethylScore follows a more coarse-grained, population-scale view compared to previous approaches [Hagmann et al. (2015)] in mainly two aspects:

- ▶ Rather than testing a large set of candidate DMRs, generated by all combinations of individual MR start and end coordinates as the attempt to find the 'natural' borders of DMRs, MethylScore considers the MR distribution of the whole population of samples. Candidate DMRs are delimited by drastic, abrupt changes in the MR frequency across all samples. The underlying assumption is that frequent changes in methylation pattern among samples are more likely 'natural' borders of DMRs.
- ▶ Rather than performing all pairwise sample comparisons on each candidate DMR, MethylScore employs a routine clustering method to assign samples into groups and solely tests groups against each other for differential methylation.

Finding candidate DMRs

For each position at which the MR frequency changes (i.e. at all MR start and end positions), a candidate DMR is defined by the following rules (illustrated in Figure 2):

- ▶ a new region is started if the MR frequency at the actual position changes by more than a user-defined percentage of the number of all samples (MethylScore parameter MR_FREQ_CHANGE) compared to the position before, or – to account for gradual small-step changes – compared to a position certain base pairs upstream (MethylScore parameter MR_FREQ_DISTANCE, default: 30 bp).
- ▶ segments need to contain a minimum number of C's covered by a minimum read coverage in at least one sample (MethylScore parameters DMR_MIN_C and DMR_MIN_COV, default: 10 and 3x, respectively).

This represents roughly a clustering of the methylated genome space into major MR frequency classes of a sample population. Longer segments could optionally be split into smaller ones using a sliding window approach.

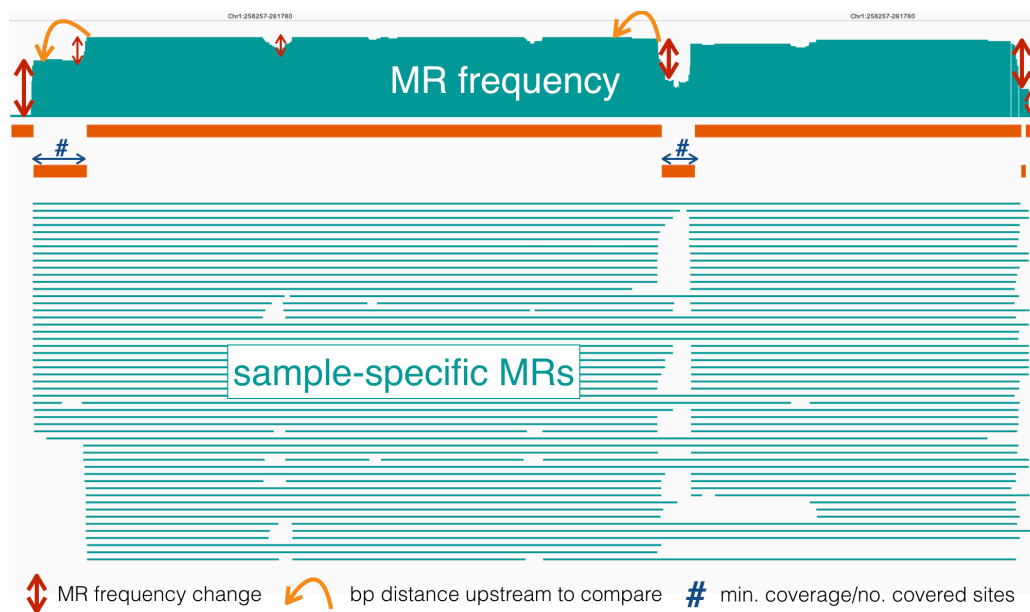


Figure 2: Determining candidate DMRs. The top panel shows an MR frequency barplot along the genome, summarizing the counts of all sample-specific MRs (the MRs are displayed as horizontal petrol-colored lines in the middle). Three parameters specify the borders of candidate DMRs (orange horizontal bars): 1. the degree of MR frequency changes along the genome (marked by red vertical arrows), 2. the upstream distance to which MR frequency is compared to (marked by orange-colored bent arrows), and 3. a minimum number and coverage of cytosines (blue arrows indicating the length of candidate DMRs).

Compared to previous methods, this approach reduces complexity, and the candidate DMRs no longer overlap each other.

Clustering and testing samples

For each candidate DMR, samples are clustered by their methylation rates using k-means clustering. Since this method requires a fixed number of clusters (k) in advance, MethylScore iteratively applies the clustering starting with $k = 2$ and subsequently incrementing by 1. In each iteration, the algorithm finds cluster centers (these centers minimize the within-cluster variances) and the methylation rate differences between all cluster centers are computed. Once a single pair of clusters differs by less than a minimum cluster methylation rate difference (MethylScore parameter `CLUSTER_MIN_METH_DIFF`, default: 10%), the iteration stops and the previous k is chosen. In most cases, a clustering into 2 groups fails to fulfill even a minimum methylation rate difference of 10 %, and hence those segments are not tested at all. If there are at least 2 clusters, these are tested against each other using the beta binomial-based statistical test that was used in Hagmann et al. (2015). The test and p-value correction (FDR calculation by Benjamini-Hochberg) is performed independently for the three sequence contexts.

The pipeline additionally assigns 'highly' differentially methylated regions (hDMRs) that exhibit 'strong' and potentially more biologically meaningful methylation differences. They include DMRs in which the mean methylation rate differs at least by a factor of 3 (default) between clusters in any context, in which the number of covered positions in a context is at least 10 (default), and in which a cluster shows at least 20% methylation (default). The latter is to prevent calling 18% methylation versus 5% methylation as 'highly differential'. This could be – especially when coverage is low – explained by false methylation rate alone, or just sampling bias.

Running MethylScore

#2.1 Prerequisites

MethylScore and its integrated helper tools are suited for Ubuntu 16.04 and rely on following installed versions of programming languages:

- ▶ perl 5.22
- ▶ python 2.7
- ▶ java 1.8

Most other required packages or modules are integrated into MethylScore. Thus, compilation is not needed. However, following python packages are required to be installed (for the python version indicated above):

- ▶ scipy
- ▶ argparse (only for stage 'igv')

#2.2 Folder structure

When MethylScore is executed, it requires a 'project folder' in which all intermediate and final files will be stored, optimally for a single project or experiment. The structure of the project folder closely follows the individual steps of the pipeline:

- ▶ **01mappings**
This folder contains all filtered (`*passQC.bam`) and deduplicated (`*passQC.dedup.bam`) mapping files. It does not contain the raw reads indicated in the sample sheet (suggested location: `00reads`).
- ▶ **02consensus**
This folder contains the output of MethylExtract per sample and chromosome (e.g. `02consensus/<sample1>/<Chr1>/allC.output`).
- ▶ **03matrix**
All genome matrices are stored in this folder. The matrix containing all analyzed samples is named `03matrix/genome_matrix.tsv`, the sample and chromosome specific ones in the respective subfolders.
- ▶ **04MRs**
This folder contains methylated regions per sample (e.g. `04MRs/<sample1>/MRs.bed`).
- ▶ **05DMRs**
Finally, here are the DMRs for the set of analyzed samples in the file `05DMRs/DMRs.bed`.
- ▶ **igv**
The igv file containing all methylation information (positional read coverage and methylation rates) will be stored in `igv/methinfo.igv`.

Each folder contains a log file (`log`) and files denoting if associated tasks have been finished (`done*` files). Only if the `done*` file is absent, the corresponding step will be performed again. To generally perform the whole pipeline and override intermediate results, set the parameter `FORCE_RERUN` to 1 in the MethylScore config file.

Logging

The standard output of the pipeline will be mirrored to the global log file of the project: `project_folder/log`. The individual stages output a `log` file for each 'job' (i.e. task of a stage) in stage-specific subfolders. When a job fails, the corresponding log file is indicated on the command line and in the global log file. When the pipeline is run on the cluster, the log files contain the peak memory usage and the runtime after the job finished. This information can be used to tweak the time and memory specifications for each stage.

#2.3 Input

MethylScore requires only an input file (sample sheet) referring to the sample and sequencing data and a configuration file containing parameter settings.

The sample sheet contains one entry per library and has to contain following columns:

1. unique sample identifier (Important: specify the same ID for technical replicates of the same sample).
2. specify "PE" if the library is paired-end, or "SE" if single-end.
3. path to the mapping file of the library's reads (produced by `bwa-meth`).
4. the (sample-specific) reference file against which the reads of that library have been mapped. To ease the specification when many/all samples share the same reference file, a specified file propagates to all following samples until the next reference file is specified. This means that in the case of a single reference file for all samples, only the first row has to contain a reference file.

Optionally, captured target regions can be provided in a ROI (regions of interest) file in standard BED format (columns: chromosome, start position, end position) via the MethylScore parameter ROI. The regions will only be used for coverage statistics, however. If analysis should only be performed on ROIs, the mapping files have to be restricted accordingly beforehand.

All parameters of the config file will be described in section #2.6.

#2.4 Output file formats

Genome matrix

The genome matrix file `03matrix/genome_matrix.tsv` contains the following columns:

1. chromosome ID
2. (1-based) position
3. sequence context (CG, CHG or CHH)
4. strand (C: forward, G: reverse)
5. ff. one string per sample: "quality/rate/meth/unmeth", where:
 - ▶ quality: is a positional quality value from MethylScore. It is listed for compatibility reasons and is not used further. It might get removed in further releases.
 - ▶ rate: is the methylation rate at this position, simply calculated as $\text{meth}/(\text{meth}+\text{unmeth})$.
 - ▶ meth/unmeth: are the numbers of methylated and unmethylated reads covering this position, respectively.

MRs

Methylated regions are provided as a BED file per sample in the folder `04MRs/`. It contains a header row for the IGV import. The columns denote:

1. chromosome ID
2. (1-based) start position
3. (1-based) end position, half-open (i.e. this position is not part of the MR)

4. number of covered cytosines in the MR
5. mean read depth of cytosines within MR
6. 75-percentile of read depth of cytosines within MR
7. mean methylation level of cytosines within MR

Here are example rows:

```
#track name=02GA0S000875_MRs color=0,153,153
Chr1    1128    2233    520     16     21     31
Chr1    2268    2711    171     41     58     21
```

DMRs

DMRs are provided as a BED file `DMRs.bed` in the folder `05DMRs/`. The columns denote:

1. chromosome ID
2. (1-based) start position
3. (1-based) end position, half-open (i.e. this position is not part of the MR)
4. Length in bp
5. Cluster String, one symbol per sample:
 - ▶ 1,2,3,... = cluster ID
 - ▶ '.' = sample is not covered at all positions within region
 - ▶ '-' = sample is not sufficiently covered at all positions within region (by default, at least 10 positions with a minimum read depth of 3-fold are required)
6. Mean methylation rate of total sites, and of sites in the contexts CG, CHG, CHH (in this order) for cluster 1, comma-separated and preceded by the cluster ID and ":" (1:...)
7. Mean methylation rate of total sites, and of sites in the contexts CG, CHG, CHH (in this order) for cluster 2, comma-separated and preceded by the cluster ID and ":" (2:...)
- 8.+ Potentially more mean methylation rates if there are more than 2 clusters

third-last Number of total sites, and of sites in the contexts CG, CHG, CHH (in this order), comma-separated, and preceded by "#:"

second-last differentially methylated contexts, comma-separated

last highly differentially methylated contexts, comma-separated

Example:

```
Chr1 1128 2233 520 11.21-211 1:15,20,0,12 2:49,80,0,28 \
1:sample1,sample2,... 2:sample6,... #:30,14,3,13 CG,CHH CG
```

The example above shows a DMR in the contexts CG and CHH (indicated in the second last column), which is at the same time a CG-hDMR (last column). Here, cluster 1 contains five samples (5 times '1' in column 5) and cluster 2 contains two samples. The cluster means in total and per context are denoted in columns 6 and 7 for cluster 1 (15 % methylation across all contexts, 20 % CG, 0 % CHG, 12 % CHH) and cluster 2, respectively. The assignment of samples to clusters follows in columns 8 and 9, and finally column 10 specifies the number of covered positions in the contexts (14 CG, 3 CHG, 13 CHH).

Across-context DMRs are also provided (`all_context_DMRs.bed`). Here, all sites across all contexts in a region have been tested. The format follows that from above without the two last columns. Instead, the last column is a flag indicating if the DMR is also an hDMR ("1") or not ("-").

Statistics

Some stages output basic statistics, which will be explained in the following.

Read statistics (dedup stage):

There is one file per sample (`01mappings/<sample>/<sample>.read_stats.tsv`) that each contains one row. To view the statistics for all samples, simply run:

```
cat 01mappings/*/*read_stats.tsv | grep -v "^#"
> 01mappings/read_stats.tsv
```

The columns are as follows:

1. **sample** sample ID
2. **total** absolute number of reads in mapping file
3. **unmapped** percentage of unmapped reads from total
4. **dupl** percentage of duplicated reads from total
5. **multpl** percentage of multiple mapping reads from total
6. **off-trg** percentage of off-target reads from total (if ROIs are provided)
7. **#ROI** absolute number of on-target reads (if ROIs are provided)
8. **ROI/flt** fraction of on-target reads from mapped, de-duplicated and uniquely mapped reads (if ROIs are provided)
9. **ROI/tot** fraction of on-target reads from total number of reads in mapping file (from column 'total'; if ROIs are provided)

Coverage statistics (dedup stage):

There is one file per sample (`01mappings/<sample>/<sample>.cov_stats.tsv`) that each contains one row. To view the statistics for all samples, simply run:

```
cat 01mappings/*/*cov_stats.tsv | grep -v "^#"
> 01mappings/cov_stats.tsv
```

The columns are as follows:

1. **sample** sample ID
2. **ON#pos** number of covered on-target sites (read depth>0)
3. **ON%pos** fraction of covered on-target sites from all sites
4. **ONavgcv** average read coverage at covered on-target sites
5. **OFF#pos** number of covered off-target sites (read depth>0)
6. **OFFavgcv** average read coverage at covered off-target sites
7. **enrich** enrichment factor (ONavgcv/OFFavgcv)

MR statistics (MRs stage):

There is one file per sample (04MRs/<sample>/MRs_stats.tsv) that each contains one row. To view the statistics for all samples, simply run:

```
cat 04MRs/*/MR_stats.tsv | grep -v "^#" > 04MRs/MR_stats.tsv
```

The columns are as follows:

1. sample ID
2. number of MRs
3. genome space in bp covered by MRs
4. average length of MRs in bp

#2.5 Performance recommendations

MethylScore can be run on a single multi-core processor or on a Sun Grid Engine-based cluster system (MethylScore parameter CLUSTER=1). The parameter THREADS correspondingly refers to the maximal number of running jobs on the cluster or to the maximal number of local cores to use.

Some steps of the MethylScore pipeline are quite I/O-intensive and parallelization might slow down runtime, especially the consensus step (MethylExtract tool). When applying MethylScore locally, we recommend to perform the consensus step on a RAM disk. In our tests, execution was nearly three times faster. When using a cluster, rather than using unlimited number of cores, it might be more efficient to limit this number (THREADS parameter) so that the file system or network will not be overloaded.

#2.6 Command line parameters

The parameters are specified in a config file, and this config file is the only minimal command line argument of MethylScore. Thus:

Quick Usage: Executing following command runs the whole pipeline on the provided test data set using the cluster, given the provided default config file:

```
MethylScore -c test/methylscore.config
```

Running MethylScore locally on a multi-core server without using a cluster can be achieved by setting CLUSTER=0 in the configuration file, or by adding the `-l` option:

```
MethylScore -c test/methylscore.config -l
```

For the execution or re-running of individual steps, provide the steps in a comma-separated list to the `-a` option. The following command is equivalent to the previous command and lists the stages that are executed by default:

```
MethylScore -c test/methylscore.config -l  
-a dedup,consensus,matrix,matrixWG,MRs,DMRs
```

The correct order of any subset of steps will be determined automatically.

MethylScore takes a few additional command line parameters, but they can be specified in the config file as well; a list of these options is printed when executing MethylScore without any argument. Command line parameters always have precedence over config file parameters.

Parameters are divided into general options and options for the individual steps. They will be listed below.

General options

Option	Default	Description
PROJECT_FOLDER	testproject/	Folder name of project, must not exist (absolute path or relative to config file location).
SAMPLE_SHEET	samplesheet.txt	White-space delimited file associating samples to their read files (absolute path or relative to config file location). For the format specification see #2.3.
ROI	ROIs.bed	File containing the enriched regions in BED format (absolute path or relative to config file location).
CLUSTER	1	Running on a SGE cluster (set to 1), or running locally on a multi-core server (set to 0).
THREADS	1	Maximal number of threads to use (when running locally), or maximal number of submitted jobs to the cluster (when CLUSTER=1).
STATISTICS	1	Flag to determine whether statistical analyses should be performed. Coverage (enrichment) analysis can take a few hours for many samples. See section #2.4.
FORCE_RERUN	0	If set to 1, all steps will be re-executed independent of whether (intermediate) results already exist. If set to 0, intermediate steps will not be performed if the <code>done</code> file of the specific step exists (cf. #2.2).
REMOVE_INTMED_FILES	1	Only keep relevant intermediate and final files.
PYTHON_PATH	python	Path of python 2.7 executable. Only needed if the system default python executable has a different version (check via <code>python --version</code>).
BAMTOOLS	bamtools	Path to bamtools executable, or "bamtools" if its binary folder is included in the <code>\$PATH</code> environment variable. Only needed if pre-compiled version of bamtools (by default in folder <code>bin_ext/</code>) is not working.

DMR options

Option	Default	Description
MR_FREQ_CHANGE	20 [%]	Minimum fraction of all samples that consistently change their (HMM-determined) methylation status at a given position. This leads to the end of the current candidate region and starts a new region here.
MR_FREQ_DISTANCE	30 [bp]	Upstream distance to which MR frequency at the focal position is compared, to check if MR_FREQ_CHANGE criterion is fulfilled.
CLUSTER_MIN_METH_DIFF	10 [%]	Minimum methylation rate difference between any pair of sample clusters. If only a single pair of clusters fails this minimum difference, the iteration of the clustering, i.e. the current k , is aborted, and $k-1$ is used as the final number of clusters.
CLUSTER_MIN_METH	20 [%]	Minimum methylation rate of any sample cluster. If this criterion is not met, the candidate region is discarded.
DMR_MIN_COV	3	Minimum read coverage of cytosines within DMRs.
DMR_MIN_C	10	Minimum number of cytosines (with at least DMR_MIN_COV coverage) within DMRs.
SLIDING_WINDOW_SIZE	0 (off)	Sliding window length in bp determining candidate DMR length.
SLIDING_WINDOW_STEP	0 (off)	Sliding window step size in bp. Must be maximally SLIDING_WINDOW_SIZE.
MR_BATCH_SIZE	1,000	Number of MR blocks per file (MR block is a genomic stretch that is methylated in at least one sample); determines the degree of parallelization of DMR calling.
HDMR_FOLD_CHANGE	3	The minimal fold change between cluster means to be called hDMR.
FDR_CUTOFF	0.05	The False Discovery Rate (FDR) cutoff for significantly differential methylation.

MR options

Option	Default	Description
MIN_COVERAGE	1	Minimum per-site read coverage. Sites with lower coverage are disregarded.
DESERT_SIZE	100 [bp]	Maximum genomic distance between adjacent covered cytosines (the region inbetween is a 'desert' since it has no methylation information). If distance exceeds this parameter, a separate HMM decoding path is started after the 'desert'. Thus, an MR cannot span a desert (see section #1.1 for further explanation).

MR_MIN_C	20	Minimum number of covered cytosines within MR. If set to 0, MRs are tested for significant length using a permutation test (see section #1.1 for details).
MERGE_DIST	30 [bp]	Distance between MRs that lead to their consolidation into a single MR.
TRIM_METHRATE	10 [%]	Maximum methylation rate of positions at each end of an MR that are trimmed off.

Consensus options

Option	Default	Description
MIN_QUAL	30	Minimum mapping quality (phred score) of reads to include in analysis. Note that the mapping quality is strongly influenced by the repetitiveness and ploidy level of the analyzed species. The default value removes a large portion of multiple mapped reads.
IGNORE_FIRST_BP	3	Chop off first bases of each read before consensus calling. This value is typically set based on M plots or on a fastQC analysis.
IGNORE_LAST_BP	1	Chop off last bases of each read before consensus calling. This value is typically set based on M plots or on a fastQC analysis.

Cluster options

There is one option for each stage that indicates how much memory (`_MEM`) and wallclock time (`_TIME`) is reserved for each job, e.g. `DEDUP_MEM` specifies the maximum RAM size for stage 1, the de-duplication of mapping files. Storage sizes are expected to be in Megabyte or Gigabyte and abbreviated by M or G, respectively. Times are specified in following format: HH:MM:SS where H: hours, M: minutes, S: seconds. Specify leading 0's, e.g. 01:00:00 to indicate 1 hour.

The runtime and peak memory the cluster job needed is printed to the log file of the respective job and can serve to tweak resource requests.

#2.7 License

The license can be found in the file `MethylScore/LICENSE` in the distributed software package.
© 2016-2018 Computomics GmbH - All rights reserved.

References

- Hagmann J, Becker C, Müller J, et al. (2015) Century-scale methylome stability in a recently diverged *arabidopsis thaliana* lineage. *PLoS genetics* **11**(1):e1004920.
- Molaro A, Hodges E, Fang F, Song Q, McCombie WR, Hannon GJ, Smith AD (2011) Sperm methylation profiles reveal features of epigenetic inheritance and evolution in primates. *Cell* **146**(6):1029–1041.
- Pedersen BS, Eyring K, De S, Yang IV, Schwartz DA (2014) Fast and accurate alignment of long bisulfite-seq reads. *arXiv preprint arXiv:1401.1129* .